

CHIEN-YI YANG

+1 (858) 241-7536 chy036@ucsd.edu  [LinkedIn](#)
 [GitHub: doctry](#)  [GitHub: doctryucsd](#)  [Portfolio](#)

Profile

Computer architecture Ph.D. candidate specializing in performance and power/energy modeling, C++/Python research tools, and hardware/software co-design for AI accelerators.

Education

University of California San Diego

Sept. 2022 – Mar. 2027 (Expected)

Ph.D. in Computer Science and Engineering

National Taiwan University

Sept. 2017 – June 2021

B.S. in Electrical Engineering

Research Experience

University of California San Diego

Sept. 2022 – Present

Graduate Research Assistant

PNMAX | Architecture and Performance Modeling

MICRO 2026 (to appear); First author

- Developed Python analytical models of compute, memory and interconnects, enabling performance projections and energy studies across HBM-PIM, UPMEM and AttAcc.
- Validated microarchitectural latency models against UniNDP, achieving **32.1× faster HBM-PIM model evaluation** with squared Pearson correlation above 0.99.
- Co-optimized accelerator architecture and tensor mappings for Llama-3, GPT-2 and OPT; used parameter sensitivity studies to identify latency/memory-footprint trade-offs and buffer, bandwidth and reduction bottlenecks.

OptiPIM | PIM Modeling and Evaluation

ISCA 2025; Distinguished Artifact Award

- Modeled and evaluated ILP-based tensor mappings, quantifying a **72.8× modeled speedup** over the Samsung HBM-PIM heuristic on the evaluated workloads.
- Validated analytical latency models against cycle-accurate Ramulator-based simulation over **12,000 mappings**, obtaining squared Pearson correlations above 0.99 for HBM-PIM and SIMDRAM.

NN-Steiner | Machine Learning for EDA

AAAI 2024

- Developed neural-algorithmic Steiner-tree optimization, reducing mean wirelength by **2.99%** versus FLUTE (A=18) across 100 uniformly sampled 5,000-point instances.
- Implemented PyTorch DistributedDataParallel/NCCL training and CPU/GPU profiling, enabling multi-GPU model updates and execution traces for training bottleneck analysis.

CSBO | Constrained Hardware/Software Optimization

IEEE ESL 2026

- Developed constraint-scheduled Bayesian optimization in BoTorch, searching a roughly 500,000-configuration design space in approximately **1.5 hours**.
- Integrated Timeloop/CiMLoop energy and latency models with PyTorch accuracy evaluation, finding CIFAR-10 trade-offs with up to **287× lower energy** or 14× lower delay than prior HDnn-PIM.

Industry Experience

MediaTek Research

Dec. 2021 – June 2022

R&D Engineer Intern

Taipei, Taiwan

Distributed T5 Training

- Deployed PyTorch Distributed with data and model parallelism, scaling T5 training to **100+ GPUs**.
- Optimized T5 training with PyTorch Profiler and critical-path analysis, achieving **30× higher throughput**.

Research Experience (continued)

University of California San Diego (continued)

Sept. 2022 – Present

Graduate Research Assistant

HD-PIM | Noise-Aware Accelerator Co-Design

ICCAD 2024

- Developed noise-aware hardware/software co-design with Bayesian optimization, identifying a MNIST design with **35.38% lower peak power** than HyperMetric at $3.4\times$ delay and 4% higher area.
- Modeled ASIC pipelines and ReRAM non-idealities, enabling power, performance and area (PPA) comparisons between digital Kronecker encoding and analog in-memory computation.

Open-Source Research

Duostra | C++ Quantum Mapping

ACM TQC 2024

- Developed C++ qubit-routing and scheduling algorithms in Duostra, integrated into qsyn, a quantum compilation framework with nearly 200 GitHub stars (Sept. 2026).
- Benchmarked graph-based mapping on circuits over 50 qubits, reducing mean mapping cost by **21.75%** versus the per-circuit best of four competing mappers.

Technical Skills

Programming: C++; Python; PyTorch; DistributedDataParallel/NCCL; PyTorch Profiler.

Architecture: Analytical performance and power/energy modeling; microarchitecture modeling; software-hardware co-design; workload characterization; bottleneck analysis; design space exploration.

Tools: Ramulator; Timeloop; CiMLoop; CACTI; BoTorch.

Selected Publications

- [1] **Chien-Yi Yang**, Jiantao Liu, Yue Pan, Tajana Rosing, Minxuan Zhou, Lana Josipović, “PNMAX: Mapping-Architecture Co-Exploration Framework for Processing Near Memory”, MICRO 2026 (to appear).
- [2] **Chien-Yi Yang**, Flavio Ponzina, Minxuan Zhou, Raid Ayoub, Pietro Mercati, Tajana Rosing, “Constraint-Scheduled Bayesian Optimization for Software-Hardware Co-Optimization on HDnn-PIM”, IEEE Embedded Systems Letters (ESL), 2026.
- [3] Esha Singh, Dongxia Wu, **Chien-Yi Yang**, Tajana Rosing, Rose Yu, Yian Ma, “Divide and Learn: Multi-Objective Combinatorial Optimization at Scale”, ICML 2026.
- [4] Jiantao Liu, Minxuan Zhou, Yue Pan, **Chien-Yi Yang**, Lana Josipović, Tajana Rosing, “OptiPIM: Optimizing Processing-in-Memory Acceleration Using Integer Linear Programming”, ISCA 2025.
- [5] **Chien-Yi Yang**, Minxuan Zhou, Flavio Ponzina, Suraj Sathya Prakash, Raid Ayoub, Pietro Mercati, Mahesh Subedar, Tajana Rosing, “Multi-Objective Software-Hardware Co-Optimization for HD-PIM via Noise-Aware Bayesian Optimization”, ICCAD 2024.
- [6] Andrew B. Kahng, Robert R. Nerem, Yusu Wang, **Chien-Yi Yang**, “NN-Steiner: A Mixed Neural-Algorithmic Approach for the Rectilinear Steiner Minimum Tree Problem”, AAAI 2024.
- [7] Chin-Yi Cheng, **Chien-Yi Yang**, Yi-Hsiang Kuo, Ren-Chu Wang, Hao-Chung Cheng, Chung-Yang (Ric) Huang, “Robust Qubit Mapping Algorithm via Double-Source Optimal Routing on Large Quantum Circuits”, ACM Transactions on Quantum Computing (TQC), 2024.